

Before you turn in the homework, make sure everything runs as expected. To do so, select **Kernel→Restart & Run All** in the toolbar above. Remember to submit both on **DataHub** and **Gradescope**.

Please fill in your name and include a list of your collaborators below.

```
In [2]: NAME = "William Sheu"  
COLLABORATORS = ""
```

Project 2: NYC Taxi Rides

Part 3: NYC Accidents Data

In the real world, data isn't always nicely bundled in one file; data can be sourced from many places with many formats. Now we will use NYC accident data to try to improve our set of features.

In this part of the project, you'll do some EDA over the combined data set. We'll do a lot of the coding work for you, but there will be a few coding subtasks for you to complete on your own, as well as many results to interpret.

Note

If your kernel dies unexpectedly, make sure you have shutdown all other notebooks. Each notebook uses valuable memory which we will need for this part of the project.

Imports

Let us start by loading the Python libraries and custom tools we will use in this part.

```
In [3]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import zipfile
import os
from pathlib import Path

sns.set(style="whitegrid", palette="muted")

plt.rcParams['figure.figsize'] = (12, 9)
plt.rcParams['font.size'] = 12

%matplotlib inline
```

Downloading the Data

We will use the `fetch_and_cache` utility to download the dataset.

```
In [4]: # Download and cache urls and get the file objects.
from utils import fetch_and_cache
data_url = 'https://github.com/DS-100/fa18/raw/gh-pages/assets/datasets'
file_name = 'collisions.zip'
dest_path = fetch_and_cache(data_url=data_url, file=file_name)

print(f'Located at {dest_path}')
```

Using version already downloaded: Fri Nov 23 03:41:41 2018
MD5 hash of file: a445b925d24f319cb60bd3ace6e4172b
Located at data/collisions.zip

We will store the taxi data locally before loading it.

```
In [5]: collisions_zip = zipfile.ZipFile(dest_path, 'r')

#Extract zip files
collisions_dir = Path('data/collisions')
collisions_zip.extractall(collisions_dir)
```

Loading and Formatting Data

The following code loads the collisions data into a Pandas DataFrame.

```
In [6]: # Run this cell to load the collisions data.
skiprows = None
collisions = pd.read_csv(collisions_dir/'collisions_2016.csv', index_col=0,
                        parse_dates={'DATETIME': ['DATE', 'TIME']}, skip
collisions['TIME'] = pd.to_datetime(collisions['DATETIME']).dt.hour
collisions['DATE'] = pd.to_datetime(collisions['DATETIME']).dt.date
collisions = collisions.dropna(subset=['LATITUDE', 'LONGITUDE'])
collisions = collisions[collisions['LATITUDE'] <= 40.85]
collisions = collisions[collisions['LATITUDE'] >= 40.63]
collisions = collisions[collisions['LONGITUDE'] <= -73.65]
collisions = collisions[collisions['LONGITUDE'] >= -74.03]
collisions.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 116691 entries, 3589202 to 3363795
Data columns (total 30 columns):
DATETIME                116691 non-null datetime64[ns]
Unnamed: 0              116691 non-null int64
BOROUGH                 100532 non-null object
ZIP CODE                100513 non-null float64
LATITUDE                116691 non-null float64
LONGITUDE               116691 non-null float64
LOCATION                 116691 non-null object
ON STREET NAME          95914 non-null object
CROSS STREET NAME       95757 non-null object
OFF STREET NAME         61545 non-null object
NUMBER OF PERSONS INJURED 116691 non-null int64
NUMBER OF PERSONS KILLED 116691 non-null int64
NUMBER OF PEDESTRIANS INJURED 116691 non-null int64
NUMBER OF PEDESTRIANS KILLED 116691 non-null int64
NUMBER OF CYCLIST INJURED 116691 non-null int64
NUMBER OF CYCLIST KILLED 116691 non-null int64
NUMBER OF MOTORIST INJURED 116691 non-null int64
NUMBER OF MOTORIST KILLED 116691 non-null int64
CONTRIBUTING FACTOR VEHICLE 1 115162 non-null object
CONTRIBUTING FACTOR VEHICLE 2 101016 non-null object
CONTRIBUTING FACTOR VEHICLE 3 7772 non-null object
CONTRIBUTING FACTOR VEHICLE 4 1829 non-null object
CONTRIBUTING FACTOR VEHICLE 5 434 non-null object
VEHICLE TYPE CODE 1       115181 non-null object
VEHICLE TYPE CODE 2       92815 non-null object
VEHICLE TYPE CODE 3       7260 non-null object
VEHICLE TYPE CODE 4       1692 non-null object
VEHICLE TYPE CODE 5       403 non-null object
TIME                      116691 non-null int64
DATE                      116691 non-null object
dtypes: datetime64[ns](1), float64(3), int64(10), object(16)
memory usage: 27.6+ MB
```

1: EDA of Accidents


Let's start by plotting the latitude and longitude where accidents occur. This may give us some insight on taxi ride durations. We sample N times (given) from the collisions dataset and create a 2D KDE plot of the longitude and latitude. We make sure to set the x and y limits according to the boundaries of New York, given below.

Here is a [map of Manhattan](#)

(<https://www.google.com/maps/place/Manhattan,+New+York,+NY/@40.7590402,-74.0394431,12z>
[73.9712488](https://www.google.com/maps/place/Manhattan,+New+York,+NY/@40.7590402,-74.0394431,12z)) for your convenience.

```
In [7]: # Plot lat/lon of accidents, will take a few seconds
N = 20000
city_long_border = (-74.03, -73.75)
city_lat_border = (40.63, 40.85)

sample = collisions.sample(N)
plt.figure(figsize=(6,8))
sns.kdeplot(sample["LONGITUDE"], sample["LATITUDE"], shade=True)
plt.xlim(city_long_border)
plt.ylim(city_lat_border)
plt.xlabel("Longitude")
plt.ylabel("Latitude")
plt.title("Accidents Distribution")
plt.show();
```



Question 1a

What can you say about the location density of NYC collisions based on the plot above?

Hint: Here is a [page](#)

(<https://www.google.com/maps/place/Manhattan,+New+York,+NY/@40.7590402,-74.0394431,73.9712488>) that may be useful, and [another page](https://www.6sqft.com/what-nycs-population-looks-like-day-vs-night/) (<https://www.6sqft.com/what-nycs-population-looks-like-day-vs-night/>) that may be useful.

```
In [8]: q1a_answer = r"""
Most of the accidents are centralllized about midtown Manhattan. This may
be because midtown has the most fluxuation of people throughout the
day, leading to more possibilities for collisions between 2 vechicles.

# YOUR CODE HERE
#raise NotImplementedError()

print(q1a_answer)
```

Most of the accidents are centralllized about midtown Manhattan. This may be because midtown has the most fluxuation of people throughout the day, leading to more possibilities for collisions between 2 vechicles.

We see that an entry in accidents contains information on number of people injured/killed. Instead of using each of these columns separately, let's combine them into one column called 'SEVERITY'. Let's also make columns FATALITY and INJURY, each aggregating the fatalities and injuries respectively.

```
In [9]: collisions['SEVERITY'] = collisions.filter(regex=r'NUMBER OF *').sum(axis=1)
collisions['FATALITY'] = collisions.filter(regex=r'KILLED').sum(axis=1)
collisions['INJURY'] = collisions.filter(regex=r'INJURED').sum(axis=1)
```

Now let's group by time and compare two aggregations: count vs mean. Below we plot the number of collisions and the mean severity of collisions by the hour, i.e. the TIME column. We visualize them side by side and set the start of our day to be 6 a.m.

Let's also take a look at the mean number of casualties per hour and the mean number of injuries per hour, plotted below.

```
In [10]: fig, axes = plt.subplots(2, 2, figsize=(16,16))
order = np.roll(np.arange(24), -6)
ax1 = axes[0,0]
ax2 = axes[0,1]
ax3 = axes[1,0]
ax4 = axes[1,1]

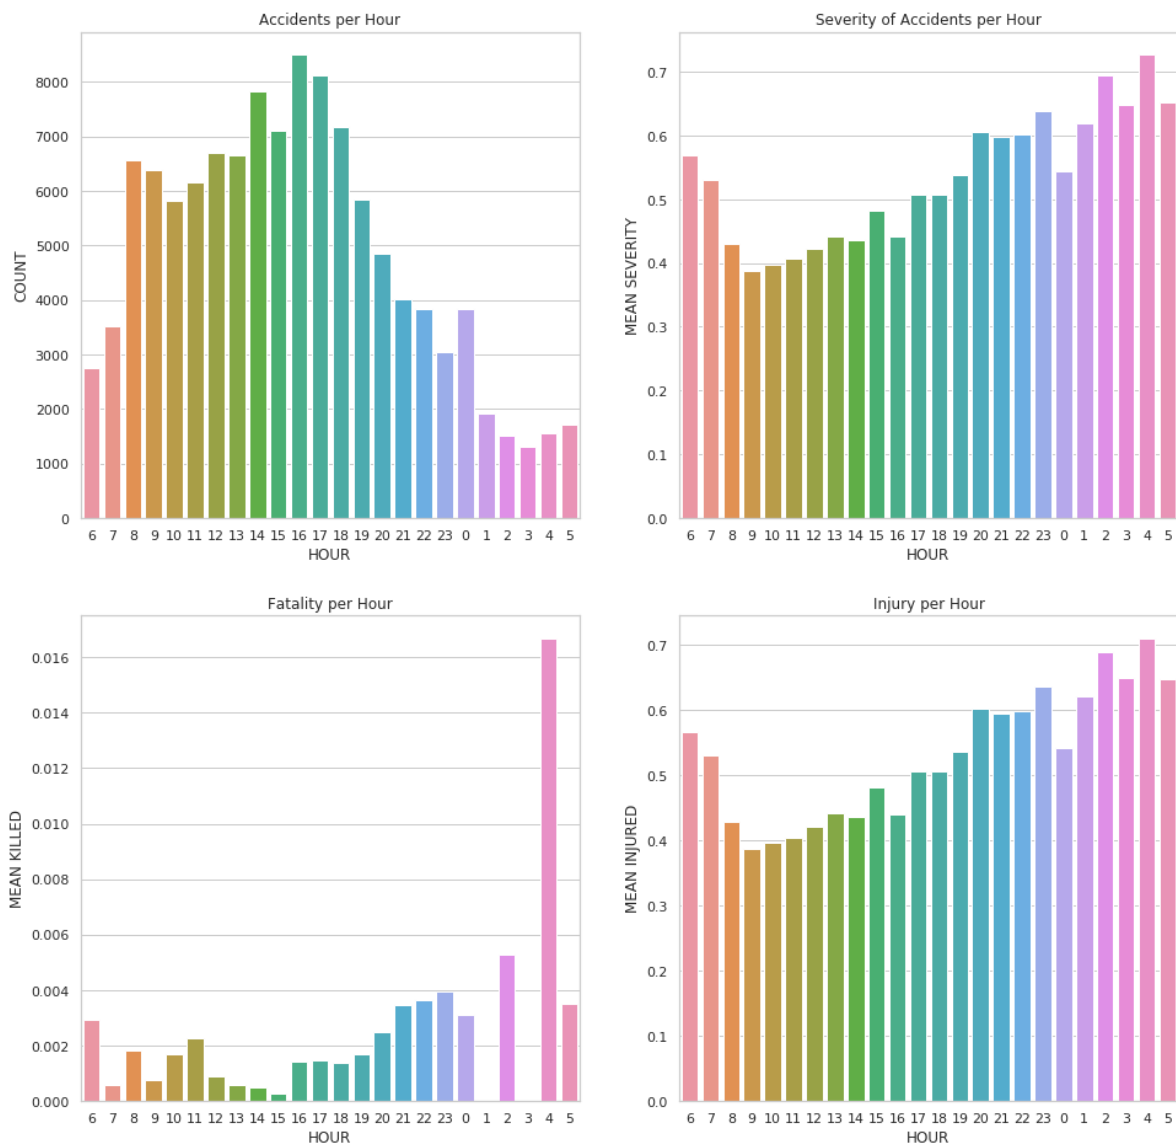
collisions_count = collisions.groupby('TIME').count()
collisions_count = collisions_count.reset_index()
sns.barplot(x='TIME', y='SEVERITY', data=collisions_count, order=order,
ax1.set_title("Accidents per Hour")
ax1.set_xlabel("HOUR")
ax1.set_ylabel('COUNT')

collisions_mean = collisions.groupby('TIME').mean()
collisions_mean = collisions_mean.reset_index()
sns.barplot(x='TIME', y='SEVERITY', data=collisions_mean, order=order,
ax2.set_title("Severity of Accidents per Hour")
ax2.set_xlabel("HOUR")
ax2.set_ylabel('MEAN SEVERITY')

fatality_count = collisions.groupby('TIME').mean()
fatality_count = fatality_count.reset_index()
sns.barplot(x='TIME', y='FATALITY', data=fatality_count, order=order, ax
ax3.set_title("Fatality per Hour")
ax3.set_xlabel("HOUR")
ax3.set_ylabel('MEAN KILLED')

injury_count = collisions.groupby('TIME').mean()
injury_count = injury_count.reset_index()
sns.barplot(x='TIME', y='INJURY', data=injury_count, order=order, ax=ax
ax4.set_title("Injury per Hour")
ax4.set_xlabel("HOUR")
ax4.set_ylabel('MEAN INJURED')

plt.show();
```



Question 1b

Based on the visualizations above, what can you say about each? Make a comparison between the accidents per hour vs the mean severity per hour. What about the number of fatalities per hour vs the number of injuries per hour? Why do we chose to have our hours start at 6 as opposed to 0?

```
In [11]: qlb_answer = r"""
```

```
The accidents per hour plot shows about what one expects: there are very  
few accidents/hour early in the morning and late in the night, when  
there are less cars about, and much more accidents/hour during the day  
time.  
The severity of accidents per hour shows that during the accidents very  
early in the morning or very late at night, the severity of the accide  
nts are on average much higher than during the daytime. This may be du  
e to the fact that there is no/less sunlight that could otherwise help  
the driver lessen the damage done in an accident. This may also be beca  
use people drive more recklessly when nobody else is on the road.  
The fatality per hour plot shows that there are less fatalities during  
the daytime, and more during the early morning/late night. This may al  
so be explained by the points above. There is also mysteriously no fat  
alities during hour 1 and hour 3, and a very sharp increase of fatalie  
s/hour during hour 4.  
The injury per hour plot basically mirrors the severity of accidents pe  
r hour plot.  
There seems to be a inverse correlation between the accidents/hour plo  
t and the severity of accidents/hour. This may be because when there i  
s more traffic, there will be more accidents, but less severe; whereas  
if there is less traffic, there are less people and thus less accident  
s, but as a result, people drive more recklessly, increasing the sever  
ity of an accident.  
The fatality per hour plot seems to be a scaled version of the injury p  
er hour plot, with less fatalies and injuries on average during dayt  
ime hours, then increase during nighttime/early morning hours. Howeve  
We chose to start our hour at 6 rather than 0 because 6 is usually seen  
"""
```

```
# YOUR CODE HERE  
#raise NotImplementedError()  
  
print(qlb_answer)
```

The accidents per hour plot shows about what one expects: there are very few accidents/hour early in the morning and late in the night, when there are less cars about, and much more accidents/hour during the day time.

The severity of accidents per hour shows that during the accidents very early in the morning or very late at night, the severity of the accidents are on average much higher than during the daytime. This may be due to the fact that there is no/less sunlight that could otherwise help the driver lessen the damage done in an accident. This may also be because people drive more recklessly when nobody else is on the road.

The fatality per hour plot shows that there are less fatalities during the daytime, and more during the early morning/late night. This may also be explained by the points above. There is also mysteriously no fatalities during hour 1 and hour 3, and a very sharp increase of fatalities/hour during hour 4.

The injury per hour plot basically mirrors the severity of accidents per hour plot.

There seems to be a inverse correlation between the accidents/hour plot and the severity of accidents/hour. This may be because when there is more traffic, there will be more accidents, but less severe; whereas if there is less traffic, there are less people and thus less accidents, but as a result, people drive more recklessly, increasing the severity of an accident.

The fatality per hour plot seems to be a scaled version of the injury per hour plot, with less fatalities and injuries on average during daytime hours, then increase during nighttime/early morning hours. However

r, there is a dramatic spike in the number of fatalities during hour 4, and there is no such spike (to that degree) during that hour in the injury per hour plot.

We chose to start our hour at 6 rather than 0 because 6 is usually seen as the start of the day, it being approximately the time of sunrise.

Let's also check the relationship between location and severity. We provide code to visualize a heat map of collisions, where the x and y coordinate are the location of the collision and the heat color is the severity of the collision. Again, we sample N points to speed up visualization.

```
In [12]: N = 10000
sample = collisions.sample(N)

# Round / bin the latitude and longitudes
sample['lat_bin'] = np.round(sample['LATITUDE'], 3)
sample['lng_bin'] = np.round(sample['LONGITUDE'], 3)

# Average severity for regions
gby_cols = ['lat_bin', 'lng_bin']

coord_stats = (sample.groupby(gby_cols)
                .agg({'SEVERITY': 'mean'})
                .reset_index())

# Visualize the average severity per region
city_long_border = (-74.03, -73.75)
city_lat_border = (40.63, 40.85)
fig, ax = plt.subplots(ncols=1, nrows=1, figsize=(14, 10))

scatter_trips = ax.scatter(sample['LONGITUDE'].values,
                           sample['LATITUDE'].values,
                           color='grey', s=1, alpha=0.5)

scatter_cmap = ax.scatter(coord_stats['lng_bin'].values,
                          coord_stats['lat_bin'].values,
                          c=coord_stats['SEVERITY'].values,
                          cmap='viridis', s=10, alpha=0.9)

cbar = fig.colorbar(scatter_cmap)
cbar.set_label("Manhattan average severity")
ax.set_xlim(city_long_border)
ax.set_ylim(city_lat_border)
ax.set_xlabel('Longitude')
ax.set_ylabel('Latitude')
plt.title('Heatmap of Manhattan average severity')
plt.axis('off');
```



Question 1c

Do you think the location of the accident has a significant impact on the severity based on the visualization above? Additionally, identify something that could be improved in the plot above and describe how we could improve it.

```
In [13]: q1c_answer = r"""  
  
No, it seems that accident location does not significantly impact the average severity of an accident.  
  
The dots on the plot seem to overlap and cover other dots, which can obscure data.  
Perhaps adding a transparency to each point will resolve this issue.  
  
"""  
  
# YOUR CODE HERE  
#raise NotImplementedError()  
  
print(q1c_answer)
```

No, it seems that accident location does not significantly impact the average severity of an accident, as the map seems to keep a relatively constant average severity.

The dots on the plot seem to overlap and cover other dots, which can obscure data. Perhaps adding a transparency to each point will resolve this issue.


Question 1d

Create a plot to visualize one or more features of the `collisions` table.

```
In [14]: # YOUR CODE HERE
collisions1=collisions.copy()
collisions1['MONTH'] = [int(str(i)[5:7]) for i in collisions['DATETIME']]

collisions1_mean = collisions1.groupby('MONTH').mean()
collisions1_mean = collisions1_mean.reset_index()
ax=sns.barplot(x='MONTH', y='SEVERITY', data=collisions1_mean)
ax.set_title("Severity of Accidents per Month")
ax.set_xlabel("MONTH")
ax.set_ylabel('MEAN SEVERITY')

plt.show();
#raise NotImplementedError()
```



Question 1e

Answer the following questions regarding your plot in 1d.

1. What feature you're visualization
2. Why you chose this feature
3. Why you chose this visualization method

```
In [15]: qle_answer = r"""  
  
I am visualizing the average severity of accidents per month. I decided  
  
"""  
# YOUR CODE HERE  
#raise NotImplementedError()  
print(qle_answer)
```

I am visualizing the average severity of accidents per month. I decided to choose this feature, thinking that some months will have more severe accidents than others, since snow and other weather conditions would greatly impact severity. I used a bar plot since it best highlights the differences between mean severity among the different months by displaying a difference in length, and since months are not a continuous distribution, I opted for a bar chart rather than a histogram.

2: Combining External Datasets

It seems like accident timing and location may influence the duration of a taxi ride. Let's start to join our NYC Taxi data with our collisions data.

Let's assume that an accident will influence traffic in the surrounding area for around 1 hour. Below, we create two columns, `START` and `END` :

- `START` : contains the recorded time of the accident
- `END` : 1 hours after `START`

Note: We chose 1 hour somewhat arbitrarily, feel free to experiment with other time intervals outside this notebook.

```
In [16]: collisions['START'] = collisions['DATETIME']  
collisions['END'] = collisions['START'] + pd.Timedelta(hours=1)
```

Question 2a

Drop all of the columns besides the following: `DATETIME` , `TIME` , `START` , `END` , `DATE` , `LATITUDE` , `LONGITUDE` , `SEVERITY` . Feel free to experiment with other subsets outside of this notebook.

```
In [17]: collisions_subset = collisions.loc[:,['DATETIME', 'TIME', 'START', 'END']
# YOUR CODE HERE
#raise NotImplementedError()
collisions_subset.head(5)
```

```
Out[17]:
```

	DATETIME	TIME	START	END	DATE	LATITUDE	LONGITUDE	SEVERITY
UNIQUE KEY								
3589202	2016-12-29 00:00:00	0	2016-12-29 00:00:00	2016-12-29 01:00:00	2016-12-29	40.844107	-73.897997	0
3587413	2016-12-26 14:30:00	14	2016-12-26 14:30:00	2016-12-26 15:30:00	2016-12-26	40.692347	-73.881778	0
3578151	2016-11-30 22:50:00	22	2016-11-30 22:50:00	2016-11-30 23:50:00	2016-11-30	40.755480	-73.741730	2
3567096	2016-11-23 20:11:00	20	2016-11-23 20:11:00	2016-11-23 21:11:00	2016-11-23	40.771122	-73.869635	0
3565211	2016-11-21 14:11:00	14	2016-11-21 14:11:00	2016-11-21 15:11:00	2016-11-21	40.828918	-73.838403	0

```
In [18]: assert collisions_subset.shape == (116691, 8)
```

Question 2b

Now, let's merge our `collisions_subset` table with `train_df`. Start by merging with only the date. We will filter by a time window in a later question.

We should be performing a left join, where our `train_df` is the left table. This is because we want to preserve all of the taxi rides in our end result. It happens that an inner join will also work, since both tables contain data on each date.

Note that the resulting `merged` table will have multiple rows for every taxi ride row in the original `train_df` table. For example, `merged` will have 483 rows with `index` equal to 16709, because there were 483 accidents that occurred on the same date as ride #16709.

Because of memory limitation, we will select the third week of 2016 to analyze. Feel free to change to it week 1 or 2 to see if the observation is general.

```
In [19]: data_file = Path("./", "cleaned_data.hdf")
train_df = pd.read_hdf(data_file, "train")
train_df = train_df.reset_index()
train_df = train_df[['index', 'tpep_pickup_datetime', 'pickup_longitude',
train_df['date'] = train_df['tpep_pickup_datetime'].dt.date
```

```
In [20]: collisions_subset = collisions_subset[collisions_subset['DATETIME'].dt.
train_df = train_df[train_df['tpep_pickup_datetime'].dt.weekofyear == 3
```

```
In [21]: # merge the dataframe here

merged = train_df.merge(collisions_subset, left_on='date', right_on='DA

# YOUR CODE HERE
#raise NotImplementedError()

merged.head()
```

```
Out[21]:
```

	index	tpep_pickup_datetime	pickup_longitude	pickup_latitude	duration	date	DATETIME	1
0	16709	2016-01-21 22:28:17	-73.997986	40.741215	736.0	2016-01-21	2016-01-21 10:35:00	
1	16709	2016-01-21 22:28:17	-73.997986	40.741215	736.0	2016-01-21	2016-01-21 13:20:00	
2	16709	2016-01-21 22:28:17	-73.997986	40.741215	736.0	2016-01-21	2016-01-21 16:00:00	
3	16709	2016-01-21 22:28:17	-73.997986	40.741215	736.0	2016-01-21	2016-01-21 18:30:00	
4	16709	2016-01-21 22:28:17	-73.997986	40.741215	736.0	2016-01-21	2016-01-21 00:05:00	

```
In [22]: assert merged.shape == (1528162, 14)
```

Question 2c

Now that our tables are merged, let's use temporal and spatial proximity to condition on the duration of the average length of a taxi ride. Let's operate under the following assumptions.

Accidents only influence the duration of a taxi ride if the following are satisfied:

- 1) The haversine distance between the the pickup location of the taxi ride and location of the recorded accident is within 5 (km). This is roughly 3.1 miles.
- 2) The start time of a taxi ride is within a 1 hour interval between the start and end of an accident.

Complete the code below to create an 'accident_close' column in the merged table that indicates if an accident was close or not according to the assumptions above.


```
In [23]: def haversine(lat1, lng1, lat2, lng2):
        """
        Compute haversine distance
        """
        lat1, lng1, lat2, lng2 = map(np.radians, (lat1, lng1, lat2, lng2))
        average_earth_radius = 6371
        lat = lat2 - lat1
        lng = lng2 - lng1
        d = np.sin(lat * 0.5) ** 2 + np.cos(lat1) * np.cos(lat2) * np.sin(l
        h = 2 * average_earth_radius * np.arcsin(np.sqrt(d))
        return h

    def manhattan_distance(lat1, lng1, lat2, lng2):
        """
        Compute Manhattan distance
        """
        a = haversine(lat1, lng1, lat1, lng2)
        b = haversine(lat1, lng1, lat2, lng1)
        return a + b
```

```
In [26]: start_to_accident = haversine(merged['pickup_latitude'].values,
                                         merged['pickup_longitude'].values,
                                         merged['LATITUDE'].values,
                                         merged['LONGITUDE'].values)
merged['start_to_accident'] = start_to_accident

# initialize accident_close column to all 0 first
merged['accident_close'] = 0

# Boolean pd.Series to select the indices for which accident_close shou
# (1) record's start_to_accident <= 5
# (2) pick up time is between start and end
is_accident_close = (merged['start_to_accident'] <= 5) & (merged['tpep_

# YOUR CODE HERE
#raise NotImplementedError()

merged.loc[is_accident_close, 'accident_close'] = 1
#for i in np.arange(len(is_accident_close)):
#    if is_accident_close[i]:
#        print(merged.iloc[[i]])
#merged[is_accident_close]
```

```
In [27]: assert merged['accident_close'].sum() > 16000
```

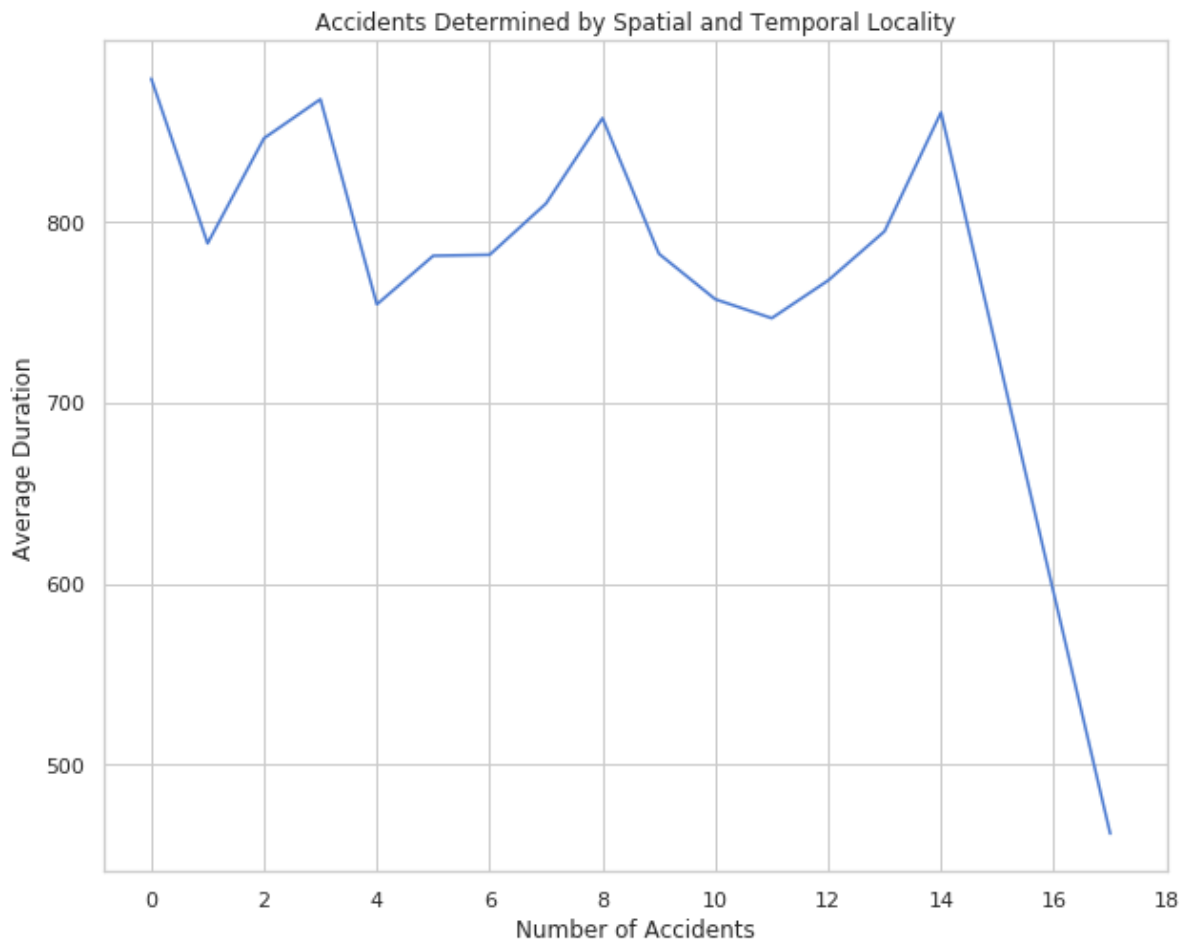
The last step is to aggregate the total number of proximal accidents. We want to count the total number of accidents that were close spatially and temporally and condition on that data.

The code below create a new data frame called `train_accidents`, which is a copy of `train_df`, but with a new column that counts the number of accidents that were close (spatially and temporally) to the pickup location/time.

```
In [28]: train_df = train_df.set_index('index')
num_accidents = merged.groupby(['index'])['accident_close'].sum().to_frame()
train_accidents = train_df.copy()
train_accidents['num_accidents'] = num_accidents
```

Next, for each value of `num_accidents`, we plot the average duration of rides with that number of accidents.

```
In [29]: plt.figure(figsize=(10,8))
train_accidents.groupby('num_accidents')['duration'].mean().plot(xticks=
plt.title("Accidents Determined by Spatial and Temporal Locality")
plt.xlabel("Number of Accidents")
plt.ylabel("Average Duration")
plt.show();
```



It seems that using both spatial and temporal proximity doesn't give us much insight on if collisions increase taxi ride durations. Let's try conditioning on spatial proximity and temporal proximity separately and see if there are more interesting results there.

```


In [30]: # Temporal locality

# Condition on time
index = (((merged['tpep_pickup_datetime'] >= merged['START']) & \
          (merged['tpep_pickup_datetime'] <= merged['END'])))

# Count accidents
merged['accident_close'] = 0
merged.loc[index, 'accident_close'] = 1
num_accidents = merged.groupby(['index'])['accident_close'].sum().to_frame()
train_accidents_temporal = train_df.copy()
train_accidents_temporal['num_accidents'] = num_accidents

# Plot
plt.figure(figsize=(10,8))
train_accidents_temporal.groupby('num_accidents')['duration'].mean().plot()
plt.title("Accidents Determined by Temporal Locality")
plt.xlabel("Number of Accidents")
plt.ylabel("Average Duration")
plt.show();

```




```
In [31]: # Spatial locality

# Condition on space
index = (merged['start_to_accident'] <= 5)

# Count accidents
merged['accident_close'] = 0
merged.loc[index, 'accident_close'] = 1
num_accidents = merged.groupby(['index'])['accident_close'].sum().to_frame()
train_accidents_spatial = train_df.copy()
train_accidents_spatial['num_accidents'] = num_accidents

# Plot
plt.figure(figsize=(10,8))
train_accidents_spatial.groupby('num_accidents')['duration'].mean().plot()
plt.title("Accidents Determined by Spatial Locality")
plt.xlabel("Number of Accidents")
plt.ylabel("Average Duration")
plt.show();
```



Question 2d

By conditioning on temporal and spatial proximity separately, we reveal different trends in average ride duration as a function of number of accidents nearby.

What can you say about the temporal and spatial proximity of accidents to taxi rides and the effect on ride duration? Think of a new hypothesis regarding accidents and taxi ride durations and explain how you would test it.

Additionally, comment on some of the assumptions being made when we condition on temporal and spatial proximity separately. What are the implications of only considering one and not the other?

```
In [32]: q2d_answer = r"""
For using just a temporal proximity of accidents, we find that that the
For using just a spatial proximity of accidents, we find that there is a
Therefore, I predict that if there are many accidents during the time of
An assumption for the temporal proximity is that if there are more acci
An assumption for the spatial proximity is that if there are less than 20
"""

# YOUR CODE HERE
#raise NotImplementedError()

print(q2d_answer)
```

For using just a temporal proximity of accidents, we find that that there is a slight positive correlation between the number of accidents of each ride and the average duration. This may be the case because if there are more accidents taking place at a certain time, then there are less roads available to drive on, leading to more traffic and thus an overall longer taxi ride.

For using just a spatial proximity of accidents, we find that there is a strong spike for the average duration if the number of accidents during the day in that area is less than ~20, then a relatively small and constant average duration for an area with 20+ accidents in the area. My theory for why this is occurring is that any taxi ride that had 20+ accidents within the vicinity of the pickup location was in Manhattan, and most trips starting at Manhattan also end in Manhattan. However, any trips with <20 accidents in the vicinity of the pickup are from the airports or the area surrounding Manhattan, where these trips should be on average travelling farther and thus longer. This might be the case, since there are bound to be more accidents within the cluttered streets of Manhattan than outside of Manhattan.

Therefore, I predict that if there are many accidents during the time of the pickup, the duration should also be longer. Also, if there are many accidents (20+) within the vicinity of the pickup during the day, then it is reasonable to assume a significantly lower duration than if there are less accidents (<20) in the area. I can test this by building a model based off of this hypothesis and off of the functions above, test it on the validation group, and seeing if my model predicts correctly or not.

An assumption for the temporal proximity is that if there are more accidents within a timeframe, there will be less roads to drive on, and thus more traffic overall. If we neglected this feature, we miss out on overall traffic conditions during the time of the taxi ride.

An assumption for the spatial proximity is that if there are less than 20 accidents that occurred in the vicinity of the pickup, the most likely

ly the pickup is in Manhattan, and vice versa. If we neglected this feature, we will not be able to approximate the location of the taxi ride and thus its average duration (unless we are clustering spatially the locations of each trip with the lat/long).

Part 3 Exports

We are not requiring you to export anything from this notebook, but you may find it useful to do so. There is a space below for you to export anything you wish.

```
In [33]: Path("data/part3").mkdir(parents=True, exist_ok=True)
data_file = Path("data/part3", "data_part3.hdf") # Path of hdf file
...
```

Out[33]: Ellipsis

```
In [34]: merged.to_hdf(data_file, "collisions")

/srv/conda/envs/data100/lib/python3.6/site-packages/pandas/core/generic.py:1996: PerformanceWarning:
your performance may suffer as PyTables will pickle object types that
it cannot
map directly to c-types [inferred_type->date,key->block3_values] [item
s->['date', 'DATE']]

return pytables.to_hdf(path_or_buf, key, self, **kwargs)
```

Part 3 Conclusions

We merged the NYC Accidents dataset with our NYC Taxi dataset, conditioning on temporal and spatial locality. We explored potential features by visualizing the relationship between number of accidents and the average duration of a ride.

Please proceed to part 4 where we will be engineering more features and building our models using a processing pipeline.

Submission

You're almost done!

Before submitting this assignment, ensure that you have:

1. Restarted the Kernel (in the menubar, select Kernel→Restart & Run All)
2. Validated the notebook by clicking the "Validate" button.

Then,

1. **Submit** the assignment via the Assignments tab in **Datahub**

2. **Upload and tag** the manually reviewed portions of the assignment on **Gradescope**

In []:

Before you turn in the homework, make sure everything runs as expected. To do so, select **Kernel→Restart & Run All** in the toolbar above. Remember to submit both on **DataHub** and **Gradescope**.

Please fill in your name and include a list of your collaborators below.

```
In [7]: NAME = "William Sheu"
        COLLABORATORS = ""
```

Project 2: NYC Taxi Rides

Part 4: Feature Engineering and Model Fitting

In this final part of the project, you will finally build a regression model that attempts to predict the duration of a taxi ride from all other available information.

You will build this model using a processing pipeline and submit your results to Kaggle. We will first walk you through a generic example using the data we saved from Part 1. Please carefully follow these steps as you will need to repeat this for your final model. After, we give you free reign and let you decide how you want to define your final model.

```
In [8]: import os
import pandas as pd
import numpy as np
import sklearn.linear_model as lm
import matplotlib.pyplot as plt
import seaborn as sns
from pathlib import Path
from sqlalchemy import create_engine
from sklearn.model_selection import cross_val_score, train_test_split, (

sns.set(style="whitegrid", palette="muted")

plt.rcParams['figure.figsize'] = (12, 9)
plt.rcParams['font.size'] = 12

%matplotlib inline
```

Training and Validation

The following code loads the training and validation data from part 1 into a Pandas DataFrame.

```
In [9]: # Run this cell to load the data.
data_file = Path("./", "cleaned_data.hdf")
train_df = pd.read_hdf(data_file, "train")
val_df = pd.read_hdf(data_file, "val")
```

Testing

Here we load our testing data on which we will evaluate your model.

```
In [10]: test_df = pd.read_csv("./proj2_test_data.csv")
test_df['tpep_pickup_datetime'] = pd.to_datetime(test_df['tpep_pickup_datetime'])
test_df.head()
```

Out[10]:

	record_id	VendorID	tpep_pickup_datetime	passenger_count	trip_distance	pickup_longitude
0	10000	1	2016-01-02 01:45:37	1	1.20	-73.982224
1	19000	2	2016-01-02 03:05:16	1	10.90	-73.999977
2	21000	1	2016-01-02 03:24:36	1	1.80	-73.986618
3	23000	2	2016-01-02 03:47:38	1	5.95	-74.002922
4	27000	1	2016-01-02 04:36:44	1	1.60	-73.986366

```
In [11]: test_df.describe()
```

Out[11]:

	record_id	VendorID	passenger_count	trip_distance	pickup_longitude	pickup_latitude
count	1.377400e+04	13774.000000	13774.000000	13774.000000	13774.000000	13774.000000
mean	3.465950e+07	1.536082	1.663642	2.954688	-72.953619	40.183693
std	2.015133e+07	0.498714	1.311739	3.704427	8.628431	4.750764
min	1.000000e+04	1.000000	0.000000	0.000000	-77.039436	0.000000
25%	1.719975e+07	1.000000	1.000000	1.000000	-73.992058	40.730610
50%	3.457400e+07	2.000000	1.000000	1.700000	-73.981846	40.759016
75%	5.216875e+07	2.000000	2.000000	3.157500	-73.967119	40.760880
max	6.940400e+07	2.000000	6.000000	104.800000	0.000000	40.861667

Modeling

We've finally gotten to a point where we can specify a simple model. Remember that we will be fitting our model on the training set we created in part 1. We will use our validation set to evaluate how well our model might perform on future data.

Reusable Pipeline

Throughout this assignment, you should notice that your data flows through a single processing pipeline several times. From a software engineering perspective, this should be sufficient motivation to abstract parts of our code into reusable functions/methods. We will now encapsulate our entire pipeline into a single function `process_data_gm`. `gm` is shorthand for "guided model".

```

In [12]: # Copied from part 2
def haversine(lat1, lng1, lat2, lng2):
    """
    Compute haversine distance
    """
    lat1, lng1, lat2, lng2 = map(np.radians, (lat1, lng1, lat2, lng2))
    average_earth_radius = 6371
    lat = lat2 - lat1
    lng = lng2 - lng1
    d = np.sin(lat * 0.5) ** 2 + np.cos(lat1) * np.cos(lat2) * np.sin(lng
    h = 2 * average_earth_radius * np.arcsin(np.sqrt(d))
    return h

# Copied from part 2
def manhattan_distance(lat1, lng1, lat2, lng2):
    """
    Compute Manhattan distance
    """
    a = haversine(lat1, lng1, lat1, lng2)
    b = haversine(lat1, lng1, lat2, lng1)
    return a + b

# Copied from part 2
def bearing(lat1, lng1, lat2, lng2):
    """
    Compute the bearing, or angle, from (lat1, lng1) to (lat2, lng2).
    A bearing of 0 refers to a NORTH orientation.
    """
    lng_delta_rad = np.radians(lng2 - lng1)
    lat1, lng1, lat2, lng2 = map(np.radians, (lat1, lng1, lat2, lng2))
    y = np.sin(lng_delta_rad) * np.cos(lat2)
    x = np.cos(lat1) * np.sin(lat2) - np.sin(lat1) * np.cos(lat2) * np.c
    return np.degrees(np.arctan2(y, x))

# Copied from part 2
def add_time_columns(df):
    """
    Add temporal features to df
    """
    df.is_copy = False # propogate write to original dataframe
    df.loc[:, 'month'] = df['tpep_pickup_datetime'].dt.month
    df.loc[:, 'week_of_year'] = df['tpep_pickup_datetime'].dt.weekofyea
    df.loc[:, 'day_of_month'] = df['tpep_pickup_datetime'].dt.day
    df.loc[:, 'day_of_week'] = df['tpep_pickup_datetime'].dt.dayofweek
    df.loc[:, 'hour'] = df['tpep_pickup_datetime'].dt.hour
    df.loc[:, 'week_hour'] = df['tpep_pickup_datetime'].dt.weekday * 24
    return df

# Copied from part 2
def add_distance_columns(df):
    """
    Add distance features to df
    """
    df.is_copy = False # propogate write to original dataframe
    df.loc[:, 'manhattan'] = manhattan_distance(lat1=df['pickup_latitude
    lng1=df['pickup_longitude

```

```

lat2=df['dropoff_latitude']
lng2=df['dropoff_longitude']

df.loc[:, 'bearing'] = bearing(lat1=df['pickup_latitude'],
                               lng1=df['pickup_longitude'],
                               lat2=df['dropoff_latitude'],
                               lng2=df['dropoff_longitude'])
df.loc[:, 'haversine'] = haversine(lat1=df['pickup_latitude'],
                                   lng1=df['pickup_longitude'],
                                   lat2=df['dropoff_latitude'],
                                   lng2=df['dropoff_longitude'])

return df

def select_columns(data, *columns):
    return data.loc[:, columns]

```

```

In [13]: def process_data_gml(data, test=False):
X = (
    data

    # Transform data
    .pipe(add_time_columns)
    .pipe(add_distance_columns)

    .pipe(select_columns,
          'pickup_longitude',
          'pickup_latitude',
          'dropoff_longitude',
          'dropoff_latitude',
          'manhattan',
          )
)
if test:
    y = None
else:
    y = data['duration']

return X, y

```

We will use our pipeline defined above to pre-process our training and test data in exactly the same way. Our functions make this relatively easy to do!

```
In [14]: # Train
X_train, y_train = process_data_gml(train_df)
X_val, y_val = process_data_gml(val_df)
guided_model_1 = lm.LinearRegression(fit_intercept=True)
guided_model_1.fit(X_train, y_train)

# Predict
y_train_pred = guided_model_1.predict(X_train)
y_val_pred = guided_model_1.predict(X_val)
```

```
/srv/conda/envs/data100/lib/python3.6/site-packages/pandas/core/generic.py:4388: FutureWarning: Attribute 'is_copy' is deprecated and will be removed in a future version.
  object.__getattr__(self, name)
/srv/conda/envs/data100/lib/python3.6/site-packages/pandas/core/generic.py:4389: FutureWarning: Attribute 'is_copy' is deprecated and will be removed in a future version.
  return object.__setattr__(self, name, value)
```

Here, `y_val` are the correct durations for each ride, and `y_val_pred` are the predicted durations based on the 7 features above (`vendorID` , `passenger_count` , `pickup_longitude` , `pickup_latitude` , `dropoff_longitude` , `dropoff_latitude` , `manhattan`).

```
In [15]: assert 600 <= np.median(y_train_pred) <= 700
assert 600 <= np.median(y_val_pred) <= 700
```

The resulting model really is a linear model just like we saw in class, i.e. the predictions are simply generated by the product $\Phi\theta$. For example, the line of code below generates a prediction for x_1 by computing $\phi_1^T \theta$. Here `guided_model_1.coef_` is θ and `X_train.iloc[0, :]` is ϕ_1 .

Note that unlike in class, here the dummy intercept term is not included in Φ .

```
In [16]: X_train.iloc[0, :].dot(guided_model_1.coef_) + guided_model_1.intercept_
```

```
Out[16]: 558.751330511368
```

We see that this prediction is exactly the same (except for possible floating point error) as generated by the `predict` function, which simply computes the product $\Phi\theta$, yielding predictions for every input.

```
In [17]: y_train_pred[0]
```

```
Out[17]: 558.75133051135344
```

In this assignment, we will use Mean Absolute Error (MAE), a.k.a. mean L1 loss, to measure the quality of our models. As a reminder, this quantity is defined as:

$$MAE = \frac{1}{n} \sum_i |y_i - \hat{y}_i|$$

Why may we want to use the MAE as a metric, as opposed to Mean Squared Error (MSE)? Using our domain knowledge that most rides are short in duration (median is roughly 600 seconds), we know that MSE is susceptible to outliers. Given that some of the outliers in our dataset are quite extreme, it is probably better to optimize for the majority of rides rather than for the outliers. You may want to remove some of these outliers later on.

```
In [18]: def mae(actual, predicted):
        """
        Calculates MAE from actual and predicted values
        Input:
            actual (1D array-like): vector of actual values
            predicted (1D array-like): vector of predicted/fitted values
        Output:
            a float, the MAE
        """

        mae = np.mean(np.abs(actual - predicted))
        return mae
```

```
In [19]: assert 200 <= mae(y_val_pred, y_val) <= 300
        print("Validation Error: ", mae(y_val_pred, y_val))
```

Validation Error: 266.136130855

Side note: scikit-learn also has tools to compute mean absolute error (`sklearn.metrics.mean_absolute_error`). In fact, most metrics that we have discussed in this class can be found as part of the [sklearn.metrics module \(https://scikit-learn.org/stable/modules/classes.html#sklearn-metrics-metrics\)](https://scikit-learn.org/stable/modules/classes.html#sklearn-metrics-metrics). Some of these may come in handy as part of your feature engineering!


Visualizing Error

You should be getting between 200 and 300 MAE, which means your model was off by roughly 3-5 minutes on trips of average length 12 minutes. This is fairly decent performance given that our basic model uses only using the pickup/dropoff latitude and manhattan distance of the trip. 3-5 minutes may seem like a lot for a trip of 12 minutes, but keep in mind that this is the *average* error. This metric is susceptible to extreme outliers, which exist in our dataset.


Now we will visualize the residual for the validation set. We will plot the following:

1. Distribution of residuals
2. Average residual grouping by ride duration

```
In [20]: # Distribution of residuals
plt.figure(figsize=(8,4))
sns.distplot(np.abs(y_val - y_val_pred))
plt.xlabel('residual')
plt.title('distribution of residuals');
```



```
In [21]: # Average residual grouping by ride duration
val_residual = X_val.copy()
val_residual['duration'] = y_val
val_residual['rounded_duration'] = np.around(y_val, -2)
val_residual['residual'] = np.abs(y_val - y_val_pred)
tmp = val_residual.groupby('rounded_duration').mean()
plt.figure(figsize=(8,4))
tmp['residual'].plot()
plt.ylabel('average residual')
plt.title('average residual by ride duration');
```



In the first visualization, we see that most of the residuals are centered around 250 seconds ~ 4 minutes. There is a minor right tail, suggesting that we are still unable to accurately fit some outliers in our data. The second visualization also suggests this, as we see the average residual

increasing as a somewhat linear function of duration. But given that our average ride duration is roughly 600-700 seconds, it seems that we are indeed optimizing for the average ride because the residuals are smallest around 600-700.

Keep this in mind when creating your final model! Visualizing the error is a powerful tool and may help diagnose shortcomings of your model. Let's go ahead and submit to kaggle, although your error on the test set may be higher than 300.

Submission to Kaggle

The following code will write your predictions on the test dataset to a CSV, which you can submit to Kaggle. You may need to modify it to suit your needs, but we recommend you make a copy and preserve the original function.

Remember that if you've performed transformations or featurization on the training data, you must also perform the same transformations on the test data in order to make predictions. For example, if you've created features for the columns `pickup_datetime` or `pickup_latitude` on the training data, you must also extract the same features in order to use scikit-learn's `.predict(...)` method.

```
In [22]: from datetime import datetime
def generate_submission(test, predictions, force=False):
    if force:
        if not os.path.isdir("submissions"):
            os.mkdir("submissions")
        submission_df = pd.DataFrame({
            "id": test_df.index.values,
            "duration": predictions,
        },
            columns=['id', 'duration'])

        timestamp = datetime.isoformat(datetime.now()).split(".")[0]

        submission_df.to_csv(f'submissions/submission_{timestamp}.csv',

        print(f'Created a CSV file: submission_{timestamp}.csv')
        print('You may now upload this CSV file to Kaggle for scoring.'
```

```
In [23]: X_test, _ = process_data_gml(test_df, True)
```

```
/srv/conda/envs/data100/lib/python3.6/site-packages/pandas/core/generi
c.py:4388: FutureWarning: Attribute 'is_copy' is deprecated and will b
e removed in a future version.
    object.__getattr__(self, name)
/srv/conda/envs/data100/lib/python3.6/site-packages/pandas/core/generi
c.py:4389: FutureWarning: Attribute 'is_copy' is deprecated and will b
e removed in a future version.
    return object.__setattr__(self, name, value)
```

```
In [24]: assert list(X_train.columns) == list(X_test.columns), "Different columns"
submission_predictions = (guided_model_1
                          .fit(X_train, y_train)
                          .predict(X_test))
submission_predictions = submission_predictions.astype(int)
submission_predictions[submission_predictions < 0] = 0
generate_submission(test_df, submission_predictions, True)
```

Created a CSV file: submission_2018-12-05T10:15:06.csv
 You may now upload this CSV file to Kaggle for scoring.

```
In [25]: # Check your submission
assert isinstance(submission_predictions, np.ndarray), "Submission not a numpy array"
assert all(submission_predictions >= 0), "Duration must be non-negative"
assert isinstance(submission_predictions.dtype.type, np.integer), "Submission predictions must be integers"
```

Your Turn!

Now it's your turn! Draw upon everything you have learned this semester to find the best features to help your model accurately predict the duration of a taxi ride.

You may use whatever method you prefer in order to create features. You may use features that we created and features that you discovered yourself from any of the 2 datasets. However, we want to make it fair to students who are seeing these techniques for the first time. As such, you are only allowed regression models and their regularized forms. This means no random forest, k-nearest-neighbors, neural nets, etc.

Here are some ideas to improve your model:

- **Data selection:** January 2016 was an odd month for taxi rides due to the blizzard. Would it help to select training data differently?
- **Data cleaning:** Try cleaning your data in different ways. In particular, consider how to handle outliers.
- **Better features:** Explore the 2 datasets and find what features are most helpful. Utilize external datasets to improve your accuracy.
- **Regularization:** Try different forms of regularization to avoid fitting to the training set. Recall that Ridge and Lasso are the names of the classes in `sklearn.linear_model` that combine `LinearRegression` with regularization techniques.
- **Model selection:** You can adjust parameters of your model (e.g., the regularization parameter) to achieve higher accuracy. [GridSearchCV \(http://scikit-learn.org/stable/modules/generated/sklearn.model_selection.GridSearchCV.html\)](http://scikit-learn.org/stable/modules/generated/sklearn.model_selection.GridSearchCV.html) may be helpful.
- **Validation:** Recall that you should use cross-validation to do feature and model selection properly! Otherwise, you will likely overfit to your training data.

There's many things you could try that could help your model. We have only suggested a few. Be creative and innovative! Please use `proj2_extras.ipynb` for all of your extraneous work. Note that you will be submitting `proj2_extras.ipynb` and we will be grading it. Please properly comment and format this notebook!

Once you are satisfied with your results, answer the questions in the Deliverables section. You may want to read this section in advance so you have an idea of what we're looking for.

Deliverables

Feature/Model Selection Process

Let's first look at selection of better features. In this following cell, describe the process of choosing good features to improve your model. You should use at least 3-4 sentences each to address the follow questions. Backup your responses with graphs supporting your claim (you can save figures and load them, no need to add the plotting code here). Use these questions to concisely summarize all of your extra work!

Question 1a

How did you find better features for your model?

```
In [26]: q1a_answer = r"""  
  
I started by looking at the train_df and seeing which columns would be useful  
in determining duration. I also tried combining multiple columns together,  
thinking that the resulting column would be useful in determining duration.  
"""  
print(q1a_answer)  
# YOUR CODE HERE  
#raise NotImplementedError()
```

I started by looking at the train_df and seeing which columns would be useful in determining duration. I also tried combining multiple columns together, thinking that the resulting column would be useful in determining duration.

Question 1b

What did you try that worked / didn't work?

```
In [27]: qlb_answer = r"""  
  
I decided to try out using fare_amount, since it would make sense for d  
  
"""  
print(qlb_answer)  
# YOUR CODE HERE  
#raise NotImplementedError()
```

I decided to try out using fare_amount, since it would make sense for duration to scale positively with the fare_amount. This worked out very well, and reduced my error significantly. Additionally, I added the hour as a feature, as at some hours, there is consistently more traffic than other hours. I also added month as a feature, since I thought the month would greatly affect duration of a ride. Both of these features slightly reduced my error.

Question 1c

What was surprising in your search for good features?

```
In [28]: qlc_answer = r"""  
  
I was expecting the fare_amount to be a great feature, but it exceeded  
  
"""  
print(qlc_answer)  
# YOUR CODE HERE  
#raise NotImplementedError()
```

I was expecting the fare_amount to be a great feature, but it exceeded my expectations, as it more than halved my error. I also added month as a feature, which slightly lowered my error, and I had to import 5 additional months worth of data.

Question 2

Just as in the guided model above, you should encapsulate as much of your workflow into functions as possible. Define `process_data_fm` and `final_model` in the cell below. In order to calculate your final model's MAE, we will run the code in the cell after that.

Note: You *MUST* name the model you wish to be evaluated on `final_model` . This is what we will be using to generate your predictions. We will take the state of `final_model` right after executing the cell below and run the following code:

```
# Load in test_df, solutions
X_test, _ = process_data_fm(test_df, True)
submission_predictions = final_model.predict(X_test)
# Generate score for autograding
```

We encourage you to conduct all of your exploratory work in `proj2_extras.ipynb` , which will be graded for 10 points.

```

In [29]: train_extra_df = pd.read_hdf(Path("data", "data_extra.hdf"), "train_ext

AVERAGE_LONG_DROP=np.mean(train_extra_df['dropoff_longitude'])
AVERAGE_LAT_DROP=np.mean(train_extra_df['dropoff_latitude'])
AVERAGE_LONG_PICK=np.mean(train_extra_df['pickup_longitude'])
AVERAGE_LAT_PICK=np.mean(train_extra_df['pickup_latitude'])

def process_data_fm(data, test=False):
    # Put your final pipeline here
    if test:
        data1 = data.copy()
    else:
        data1 = data.copy()[ (data['duration'] < 10000) ]
    if test:
        data1.loc[(data1['dropoff_longitude'] > -70), 'dropoff_longitude'] =
        data1.loc[(data1['dropoff_latitude'] < 35), 'dropoff_latitude'] =
        data1.loc[(data1['pickup_longitude'] > -70), 'pickup_longitude'] =
        data1.loc[(data1['pickup_latitude'] < 35), 'pickup_latitude'] =
    X = (data1.pipe(add_time_columns).pipe(add_distance_columns)
        .pipe(select_columns,
              'pickup_longitude',
              'pickup_latitude',
              'dropoff_longitude',
              'dropoff_latitude',
              'manhattan',
              'fare_amount',
              'hour',
              'month'
              )
        )
    if test:
        y = None
    else:
        y = data1['duration']

    return X, y

X_train2, y_train2 = process_data_fm(train_extra_df)
final_model = lm.LinearRegression(fit_intercept=True)
final_model.fit(X_train2, y_train2)

# YOUR CODE HERE
#raise NotImplementedError()

```

```

/srv/conda/envs/data100/lib/python3.6/site-packages/pandas/core/generi
c.py:4388: FutureWarning: Attribute 'is_copy' is deprecated and will b
e removed in a future version.
    object.__getattr__(self, name)
/srv/conda/envs/data100/lib/python3.6/site-packages/pandas/core/generi
c.py:4389: FutureWarning: Attribute 'is_copy' is deprecated and will b
e removed in a future version.
    return object.__setattr__(self, name, value)

```

```

Out[29]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=1, normalize=
False)

```

```
In [30]: # Feel free to change this cell
X_test, _ = process_data_fm(test_df, True)
final_predictions = final_model.predict(X_test)
final_predictions = final_predictions.astype(int)
generate_submission(test_df, final_predictions, False) # Change to true
print(final_predictions)
```

```
[ 409 1861  544 ...,  964  632  437]
```

```
/srv/conda/envs/data100/lib/python3.6/site-packages/pandas/core/generi
c.py:4388: FutureWarning: Attribute 'is_copy' is deprecated and will b
e removed in a future version.
  object.__getattr__(self, name)
/srv/conda/envs/data100/lib/python3.6/site-packages/pandas/core/generi
c.py:4389: FutureWarning: Attribute 'is_copy' is deprecated and will b
e removed in a future version.
  return object.__setattr__(self, name, value)
```

Question 3

The following hidden cells will test your model on the test set. Please do not delete any of them if you want credit!

```
In [31]: # NO TOUCH
```

```
In [32]: # NOH
```

```
In [33]: # STAHP
```

```
In [34]: # NO MOLESTE
```

```
In [35]: # VA-T'EN
```

```
In [36]: # NEIN
```

```
In [37]: # PLSNO
```

```
In [38]: # THIS SPACE IS NOT YOURS
```

```
In [39]: # TAWDEETAW
```

```
In [40]: # MAU LEN
```

```
In [41]: # ALMOST
```

```
In [42]: # TO
```

```
In [43]: # THE
```

```
In [44]: # END
```

```
In [45]: # Hmph
```

```
In [46]: # Good riddance
```

```
In [47]: generate_submission(test_df, submission_predictions, True)
```

Created a CSV file: submission_2018-12-05T10:15:07.csv
You may now upload this CSV file to Kaggle for scoring.

This should be the format of your CSV file.

Unix-users can verify it running `!head submission_{datetime}.csv` in a jupyter notebook cell.

```
id,duration
id3004672,965.3950873305439
id3505355,1375.0665915134596
id1217141,963.2285454171943
id2150126,1134.7680929570924
id1598245,878.5495792656438
id0668992,831.6700312449248
id1765014,993.1692116960185
id0898117,1091.1171629594755
id3905224,887.9037911118357
```

Kaggle link: <https://www.kaggle.com/t/f8b3c6acc3a045cab152060a5bc79670>
(<https://www.kaggle.com/t/f8b3c6acc3a045cab152060a5bc79670>)

Submission

You're almost done!

Before submitting this assignment, ensure that you have:

1. Restarted the Kernel (in the menubar, select Kernel→Restart & Run All)
2. Validated the notebook by clicking the "Validate" button.

Then,

1. **Submit** the assignment via the Assignments tab in **Datahub**
2. **Upload and tag** the manually reviewed portions of the assignment on **Gradescope**

```
In [ ]:
```


Before you turn in the homework, make sure everything runs as expected. To do so, select **Kernel→Restart & Run All** in the toolbar above. Remember to submit both on **DataHub** and **Gradescope**.

Please fill in your name and include a list of your collaborators below.

```
In [1]: NAME = "William Sheu"
        COLLABORATORS = ""
```

Project 2: NYC Taxi Rides

Extras

Put all of your extra work in here. Feel free to save figures to use when completing Part 4.

```
In [2]: import os
import pandas as pd
import numpy as np
import sklearn.linear_model as lm
import matplotlib.pyplot as plt
import seaborn as sns
import zipfile
from utils import fetch_and_cache
from pathlib import Path
from sqlalchemy import create_engine
from datetime import datetime
from sklearn.model_selection import cross_val_score, train_test_split, (
```

```

In [3]: test_df = pd.read_csv("./proj2_test_data.csv")
DB_URI = "sqlite:///srv/db/taxi_2016_student_small.sqlite"
TABLE_NAME = "taxi"
query = """
    SELECT *
    FROM (
        SELECT *
        FROM (
            SELECT *, julianday(tpep_dropoff_datetime) - julianday(tpep_pickup_datetime) AS duration
            FROM (
                SELECT *
                FROM taxi
                WHERE tpep_pickup_datetime
                    BETWEEN '2016-01-01' AND '2016-07-01'
                    AND record_id % 100 == 0
                ORDER BY tpep_pickup_datetime
            )
        )
    WHERE duration < 0.1157407
    )
    WHERE (
        pickup_longitude <= -73.75 AND
        pickup_longitude >= -74.03 AND
        dropoff_longitude <= -73.75 AND
        dropoff_longitude >= -74.03 AND
        pickup_latitude <= 40.85 AND
        pickup_latitude >= 40.63 AND
        dropoff_latitude <= 40.85 AND
        dropoff_latitude >= 40.63
    )
    WHERE (passenger_count > 0)"""
sql_engine = create_engine(DB_URI)
processed_df = pd.read_sql_query(query, sql_engine)
processed_df['tpep_pickup_datetime'] = pd.to_datetime(processed_df['tpep_pickup_datetime'])
processed_df['tpep_dropoff_datetime'] = pd.to_datetime(processed_df['tpep_dropoff_datetime'])
processed_df['duration'] = processed_df['duration']*86400

```

```

In [4]: # Copied from part 2
def haversine(lat1, lng1, lat2, lng2):
    """
    Compute haversine distance
    """
    lat1, lng1, lat2, lng2 = map(np.radians, (lat1, lng1, lat2, lng2))
    average_earth_radius = 6371
    lat = lat2 - lat1
    lng = lng2 - lng1
    d = np.sin(lat * 0.5) ** 2 + np.cos(lat1) * np.cos(lat2) * np.sin(l
    h = 2 * average_earth_radius * np.arcsin(np.sqrt(d))
    return h

# Copied from part 2
def manhattan_distance(lat1, lng1, lat2, lng2):
    """
    Compute Manhattan distance
    """
    a = haversine(lat1, lng1, lat1, lng2)
    b = haversine(lat1, lng1, lat2, lng1)
    return a + b

# Copied from part 2
def bearing(lat1, lng1, lat2, lng2):
    """
    Compute the bearing, or angle, from (lat1, lng1) to (lat2, lng2).
    A bearing of 0 refers to a NORTH orientation.
    """
    lng_delta_rad = np.radians(lng2 - lng1)
    lat1, lng1, lat2, lng2 = map(np.radians, (lat1, lng1, lat2, lng2))
    y = np.sin(lng_delta_rad) * np.cos(lat2)
    x = np.cos(lat1) * np.sin(lat2) - np.sin(lat1) * np.cos(lat2) * np.c
    return np.degrees(np.arctan2(y, x))

# Copied from part 2
def add_time_columns(df):
    """
    Add temporal features to df
    """
    df.is_copy = False # propogate write to original dataframe
    df.loc[:, 'month'] = df['tpep_pickup_datetime'].dt.month
    df.loc[:, 'week_of_year'] = df['tpep_pickup_datetime'].dt.weekofyea
    df.loc[:, 'day_of_month'] = df['tpep_pickup_datetime'].dt.day
    df.loc[:, 'day_of_week'] = df['tpep_pickup_datetime'].dt.dayofweek
    df.loc[:, 'hour'] = df['tpep_pickup_datetime'].dt.hour
    df.loc[:, 'week_hour'] = df['tpep_pickup_datetime'].dt.weekday * 24
    return df

# Copied from part 2
def add_distance_columns(df):
    """
    Add distance features to df
    """
    df.is_copy = False # propogate write to original dataframe
    df.loc[:, 'manhattan'] = manhattan_distance(lat1=df['pickup_latitude
    lng1=df['pickup_longitude

```

```

lat2=df['dropoff_latitude']
lng2=df['dropoff_longitude']

df.loc[:, 'bearing'] = bearing(lat1=df['pickup_latitude'],
                               lng1=df['pickup_longitude'],
                               lat2=df['dropoff_latitude'],
                               lng2=df['dropoff_longitude'])
df.loc[:, 'haversine'] = haversine(lat1=df['pickup_latitude'],
                                   lng1=df['pickup_longitude'],
                                   lat2=df['dropoff_latitude'],
                                   lng2=df['dropoff_longitude'])

return df

def select_columns(data, *columns):
    return data.loc[:, columns]

def mae(actual, predicted):
    """
    Calculates MAE from actual and predicted values
    Input:
        actual (1D array-like): vector of actual values
        predicted (1D array-like): vector of predicted/fitted values
    Output:
        a float, the MAE
    """
    mae = np.mean(np.abs(actual - predicted))
    return mae

def generate_submission(test, predictions, force=False):
    if force:
        if not os.path.isdir("submissions"):
            os.mkdir("submissions")
        submission_df = pd.DataFrame({
            "id": test_df.index.values,
            "duration": predictions,
        },
            columns=['id', 'duration'])

        timestamp = datetime.isoformat(datetime.now()).split(".")[0]

        submission_df.to_csv(f'submissions/submission_{timestamp}.csv',

        print(f'Created a CSV file: submission_{timestamp}.csv')
        print('You may now upload this CSV file to Kaggle for scoring.')

```

```

In [5]: train_df, val_df = train_test_split(processed_df, test_size=0.2, random

AVERAGE_LONG_DROP=np.mean(train_df['dropoff_longitude'])
AVERAGE_LAT_DROP=np.mean(train_df['dropoff_latitude'])
AVERAGE_LONG_PICK=np.mean(train_df['pickup_longitude'])
AVERAGE_LAT_PICK=np.mean(train_df['pickup_latitude'])

def process_data_gm2(data, test=False):
    if test:
        data1 = data.copy()
    else:
        data1 = data.copy()[data['duration'] < 10000]
    if test:
        data1.loc[(data1['dropoff_longitude'] > -70), 'dropoff_longitude']
        data1.loc[(data1['dropoff_latitude'] < 35), 'dropoff_latitude']
        data1.loc[(data1['pickup_longitude'] > -70), 'pickup_longitude']
        data1.loc[(data1['pickup_latitude'] < 35), 'pickup_latitude'] =
X = (data1.pipe(add_time_columns).pipe(add_distance_columns)
     .pipe(select_columns,
           'pickup_longitude',
           'pickup_latitude',
           'dropoff_longitude',
           'dropoff_latitude',
           'manhattan',
           'fare_amount',
           'hour',
           'month'
          )
    )
    if test:
        y = None
    else:
        y = data1['duration']

    return X, y

```

```

In [6]: X_train, y_train = process_data_gm2(train_df)
X_val, y_val = process_data_gm2(val_df)
guided_model_2 = lm.LinearRegression(fit_intercept=True)
guided_model_2.fit(X_train, y_train)
y_val_pred = guided_model_2.predict(X_val)
print(mae(y_val_pred, y_val))

```

```

/srv/conda/envs/data100/lib/python3.6/site-packages/pandas/core/generi
c.py:4388: FutureWarning: Attribute 'is_copy' is deprecated and will b
e removed in a future version.

```

```

    object.__getattr__(self, name)

```

```

/srv/conda/envs/data100/lib/python3.6/site-packages/pandas/core/generi
c.py:4389: FutureWarning: Attribute 'is_copy' is deprecated and will b
e removed in a future version.

```

```

    return object.__setattr__(self, name, value)

```

```

175.869207619

```

```
In [7]: print(guided_model_2.coef_)
train_df.head(10)
```

```
[-1466.7188094    923.73865708 -1740.60100377   107.43336519    2.502
18244
   61.37640329    2.88886216    10.32365152]
```

```
Out[7]:
```

	record_id	VendorID	tpep_pickup_datetime	tpep_dropoff_datetime	passenger_count	trip
30390	15892900	1	2016-02-10 14:46:02	2016-02-10 15:13:52	1	
124167	59546000	2	2016-06-03 22:19:02	2016-06-03 22:25:05	1	
65114	29985300	2	2016-03-23 10:50:48	2016-03-23 10:59:22	1	
113090	54223000	2	2016-05-20 13:01:23	2016-05-20 13:08:20	1	
92178	44175100	1	2016-04-25 08:39:56	2016-04-25 08:42:36	1	
43260	19244300	1	2016-02-26 05:53:23	2016-02-26 05:55:53	1	
98668	47420300	1	2016-05-03 08:53:52	2016-05-03 09:05:52	1	
61205	27975300	1	2016-03-18 14:01:48	2016-03-18 14:09:26	1	
135630	64839900	1	2016-06-18 11:17:10	2016-06-18 11:34:00	1	
91719	43943300	1	2016-04-24 14:09:57	2016-04-24 14:28:49	3	

10 rows × 21 columns

```
In [8]: test_df['tpep_pickup_datetime'] = pd.to_datetime(test_df['tpep_pickup_datetime'])
test_df = test_df.pipe(add_distance_columns)
test_df[test_df['manhattan'] == 0].loc[:, ['pickup_longitude', 'pickup_latitude']]
```

```
/srv/conda/envs/data100/lib/python3.6/site-packages/pandas/core/generic.py:4388: FutureWarning: Attribute 'is_copy' is deprecated and will be removed in a future version.
  object.__getattr__(self, name)
/srv/conda/envs/data100/lib/python3.6/site-packages/pandas/core/generic.py:4389: FutureWarning: Attribute 'is_copy' is deprecated and will be removed in a future version.
  return object.__setattr__(self, name, value)
```

```
In [10]: X_test, _ = process_data_gm2(test_df, True)
final_predictions = guided_model_2.predict(X_test)
final_predictions = final_predictions.astype(int)
generate_submission(test_df, final_predictions, False) # Change to true
final_predictions
```

Created a CSV file: submission_2018-12-04T23:52:45.csv
You may now upload this CSV file to Kaggle for scoring.

```
/srv/conda/envs/data100/lib/python3.6/site-packages/pandas/core/generi
c.py:4388: FutureWarning: Attribute 'is_copy' is deprecated and will b
e removed in a future version.
  object.__getattr__(self, name)
/srv/conda/envs/data100/lib/python3.6/site-packages/pandas/core/generi
c.py:4389: FutureWarning: Attribute 'is_copy' is deprecated and will b
e removed in a future version.
  return object.__setattr__(self, name, value)
```

```
Out[10]: array([ 409, 1861,  544, ...,  964,  632,  437])
```

```
In [12]: data_file = Path("data", "data_extra.hdf") # Path of hdf file
train_df.to_hdf(data_file, "train_extra_df")
```

Submission

You're almost done!

Before submitting this assignment, ensure that you have:

1. Restarted the Kernel (in the menubar, select Kernel→Restart & Run All)
2. Validated the notebook by clicking the "Validate" button.

Then,

1. **Submit** the assignment via the Assignments tab in **Datahub**
2. **Upload and tag** the manually reviewed portions of the assignment on **Gradescope**